

Sentiment Analysis of Amazon Reviews using NLP

I started this project by cleaning Amazon product review data using Python libraries like `re` & `string`, removing unwanted characters, numbers, symbols, and emojis. After preprocessing, I converted all text to lowercase to keep it consistent.

I then created additional features like **Review Length**, **Word Count**, **Exclamation Marks Count**, & **Sentiment Score** using `TextBlob`, helping the model better understand the reviews. For vectorizing the cleaned text, I used `TfidfVectorizer` from scikit-learn.

I trained multiple models including **Logistic Regression**, **Random Forest**, **Neural Network**, **Naive Bayes**, **SVM**, **Voting Classifier**, & **LSTM**, & evaluated them using cross-validated **F1 Scores**. Based on the comparison, the **Neural Network** gave the best performance. Using **automatic logic**, the best-performing model, along with the **TF-IDF vectorizer** & **scaler**, were saved using `joblib` for future predictions without needing retraining.

In addition, I used **Hugging Face's bertweet-base-sentiment-analysis** model through a pipeline for side-by-side sentiment predictions in the app as a reference, but it was not included in the F1 score-based evaluation or model selection process.

To make the solution interactive, I built a **Streamlit dashboard** where users can type their own review or quickly select predefined examples (Positive, Neutral, Negative) using '**Try an Example**' buttons. After clicking **Predict**, the app shows the predicted sentiment along with a **Review Analysis** section displaying **Review Length**, **Word Count**, **Exclamation Marks**, **Emoji Count**, & **Sentiment Score**. A **Confidence Breakdown Pie Chart** visually represents the model's prediction probabilities, and users can easily reset all fields or download the prediction result as a **CSV file**.



by Debasis Baidya



Contributions of Key Libraries & Modules to the Project

1

Pandas & Numpy

I used these for loading, cleaning, and manipulating structured data and handling numerical computations.



2

re, NLTK & TextBlob

I used these for cleaning tweet text, removing stopwords, lemmatizing, and calculating sentiment polarity.



3

Matplotlib, Seaborn & WordCloud

I used these for visualizing EDA insights like class distribution, word frequencies, and generating word clouds.



4

Scikit-learn (sklearn)

I used this for TF-IDF vectorization, building ML models, evaluating performance using accuracy, F1-score, & confusion matrix.



5

Scipy.sparse

I used this for combining sparse TF-IDF features with engineered features like sentiment score, tweet length, and hashtag count.



6

Joblib & Pickle

I used these for saving and loading trained models and vectorizers for real-time prediction.



7

Num2words, contractions, emoji

I used num2words to convert numbers to words, contractions to expand shortened words, & emoji to count & analyze in the input text.



8

OS & IPython.display

I used these for handling files/directories & improving output readability in the notebook.



9

Streamlit

I used this to build and deploy an interactive web app for real-time tweet classification.



10

Transformers (Hugging Face)

I used this optionally to compare traditional models with state-of-the-art deep learning models.



```

1 # Handling Missing Values of Train Data
2 print("Missing Values Before Handling:")
3 print(train_df.isnull().sum()[train_df.isnull().sum() > 0])
4
5 train_df.fillna("", inplace=True) # Filling missing values with an empty string
6
7 print("\nMissing Values After Handling:")
8 print(train_df.isnull().sum())

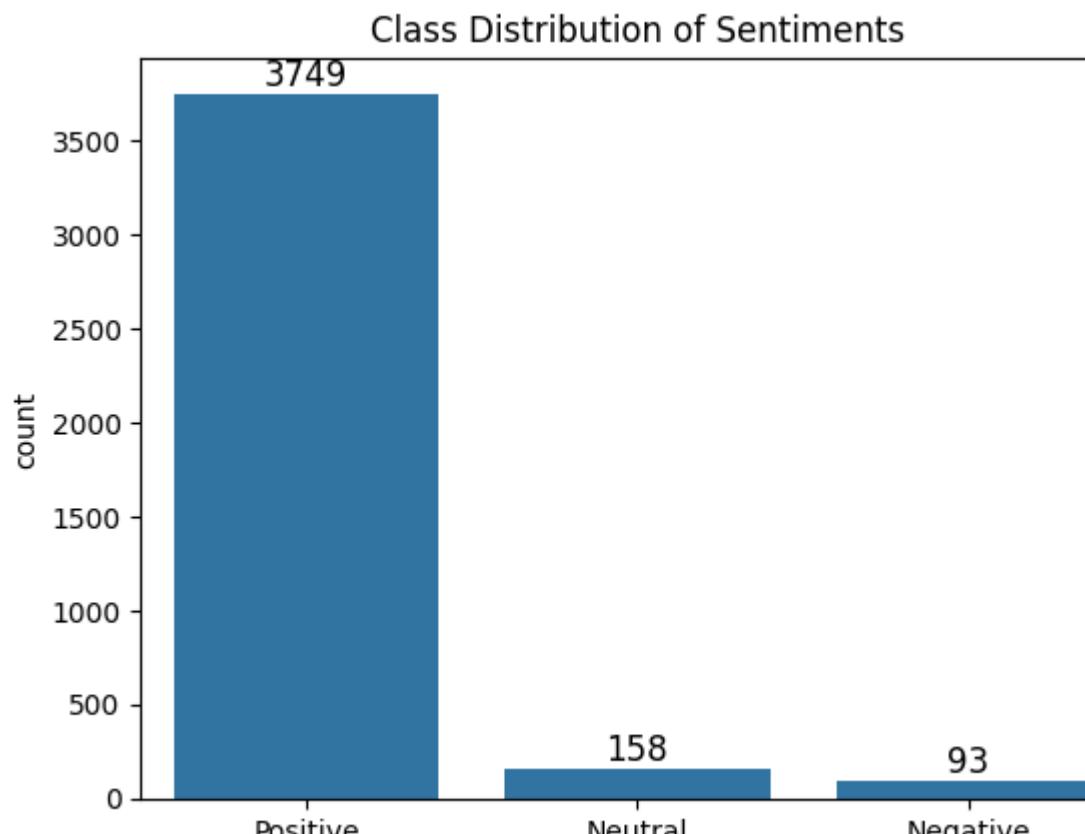
```

Missing Values Before Handling:

```
reviews.title    10
dtype: int64
```

Missing Values After Handling:

```
Name of the product      0
Product Brand            0
categories                0
primaryCategories        0
reviews.date              0
reviews.text              0
reviews.title              0
sentiment                  0
dtype: int64
```



Dataset Overview

📁 Data Loading

Loaded dataset with:

- ❑ **train_data.csv**: 4000 rows, 8 columns
- ❑ **test_data.csv**: 1000 rows, 7 columns

Includes product info, reviews, and sentiment (train only).

📊 Dataset Overview

- ❑ Train Data has an extra column: sentiment (target variable).
- ❑ Data Types (all columns are object type - text).

⚠️ Handling Missing Values

- ❑ Train Data: reviews.title - 10 missing values.
- ❑ Test Data: reviews.title - 3 missing values.
- Handled by replacing with empty string (.fillna("")).

📊 Class Distribution of Sentiments

- ❑ Positive: 3749
- ❑ Neutral: 158
- ❑ Negative: 93

📈 Product Categories

- ❑ Popular categories include Electronics, Home & Kitchen, & Clothing.

Label Encoding, TF-IDF, Feature Engineering & Tackling Imbalance

1 Label Encoding of Sentiment

1

Encoded sentiment labels:

- 'Negative' → 0, 'Neutral' → 1, 'Positive' → 2

Saved encoder as label_encoder.pkl for consistency

2 TF-IDF Vectorization

2

max_features=8000, ngram_range=(1, 2)

min_df=3, max_df=0.9, sublinear_tf=True for optimization

Saved vectorizer as vectorizer.pkl for reuse.

3 Extracted review-level features

3

text_length (characters) | word_count (words) | ! exclamtion_count

Combined with TF-IDF vectors for enriched features.

4 Tackling Class Imbalance

4

Applied SMOTETomek to balance class distribution.

X Before: ☺ Positive: 3749, ☻ Neutral: 158, ☹ Negative: 93

✓ After: ☺ Positive: 3738, ☻ Neutral: 3732, ☹ Negative: 3741

5 Saved Preprocessing Files

5

Saved key preprocessing objects for future use:

- label_encoder.pkl (for consistent encoding)
- vectorizer.pkl (for TF-IDF transformations)

Ensures reproducibility and smooth workflow.

6 Summary & Readiness

6

✍ Data is cleaned, enriched with extra features.

⚖️ Balanced for training, ready for sentiment classifier.

```
1 # =====
2 # 3. Preprocessing & Text Normalization (Optimized)
3 # =====
4
5 import re # For regex-based cleaning
6 import nltk # Natural Language Toolkit
7 import string # For punctuation
8 import contractions # handles all common contractions
9 from num2words import num2words # handles ordinals
10
11 # Download necessary resources from NLTK
12 nltk.download('stopwords')
13 nltk.download('wordnet')
14 nltk.download(['punkt', 'punkt_tab'])
15 nltk.download('omw-1.4') # for better lemmatization
16
17 # Import after downloading
18 from nltk.corpus import stopwords
19 from nltk.stem import WordNetLemmatizer
20
21 # Initialize tools
22 lemmatizer = WordNetLemmatizer()
23 stop_words = set(stopwords.words('english'))
24
25 # Expand ordinals like "1st", "2nd", etc. → "first", "second", etc.
26 def convert_ordinals(text):
27     return re.sub(r'\b(\d+)(st|nd|rd|th)\b',
28                  lambda m: num2words(int(m.group(1)), to='ordinal'),
29                  text)
30
31 def clean_text(text):
32     text = str(text).lower() # lowercase everything
33     text = convert_ordinals(text) # convert ordinals to words
34     text = contractions.fix(text) # expand contractions
35
36     # Remove unwanted patterns: URLs, digits, mentions, punctuations, extra spaces
37     text = re.sub(r"http\S+", "", text) # remove URLs
38     text = re.sub(r"\d+", "", text) # remove digits
39     text = re.sub(r"@[\w\W]+", "", text) # remove mentions
40     text = re.sub(r"[\w\W]+?", "", text) # remove punctuations
41     text = re.sub(r"\s+", " ", text).strip() # normalize extra spaces
42
43     # Reduce repeated characters (e.g., "soooo" → "soo")
44     text = re.sub(r"(.)(\1{2,})", r"\1\1", text)
45
46     # Tokenize, remove stopwords, and lemmatize
47     words = [
48         lemmatizer.lemmatize(word)
49         for word in text.split()
50         if word not in stop_words and (len(word) > 1 or word in ['no', 'ok', 'go'])
51     ]
52
53     return " ".join(words)
54
55 # Apply preprocessing on both train and test datasets
56 train_df['cleaned_text'] = train_df['reviews.text'].astype(str).apply(clean_text)
57 test_df['cleaned_text'] = test_df['reviews.text'].astype(str).apply(clean_text)
58
59 # Expanded neutral keyword list for rule-based checks
60 neutral_keywords = [
61     'okay', 'fine', 'average', 'meh', 'just okay', 'not that much', 'not bad',
62     'mediocre', 'so-so', 'alright', 'nothing special', 'kind of', 'could be better',
63     'couldn\'t care less', 'indifferent', 'okay-ish', 'neither good nor bad',
64     'passable', 'acceptable', 'not great', 'nothing remarkable', 'alright-ish',
65     'just fine', 'could be worse', 'not bad, not good', 'somewhat okay', 'meh, could be better',
66     'nothing to complain about', 'barely noticeable', 'average at best', 'mediocre at best', 'tolerable'
67 ]
68
69 # Add neutral keyword flag for rule-based checks
70 train_df['has_neutral_keyword'] = train_df['reviews.text'].str.lower().apply(
71     lambda x: any(kw in x for kw in neutral_keywords)
72 )
73 test_df['has_neutral_keyword'] = test_df['reviews.text'].str.lower().apply(
74     lambda x: any(kw in x for kw in neutral_keywords)
75 )
76
77 # Show sample cleaned text
78 print("\nSample cleaned text:")
79 print(train_df[['reviews.text', 'cleaned_text']].head())
```

```
1 # Fit the encoder on the training set and transform the 'Sentiment' column to numeric labels
2 train_df['sentiment'] = le.fit_transform(train_df['sentiment'])
3
4 # Print the mapping of labels to their corresponding classes
5 print("\nLabel mapping:", dict(zip(le.classes_, range(len(le.classes_)))))
6
7 # Define the filename for saving the LabelEncoder
8 encoder_filename = 'label_encoder.pkl'
9
10 # Save the LabelEncoder to a .pkl file for future use
11 import joblib
12 with open(encoder_filename, 'wb') as file:
13     joblib.dump(le, file)
14
15 # Display a message indicating the encoder has been saved
16 print(f"\nLabel encoding setup complete and '{encoder_filename}' saved.")

Label mapping: {'Negative': 0, 'Neutral': 1, 'Positive': 2}
Label encoding setup complete and 'label_encoder.pkl' saved.
```

```
1 # =====
2 # 5. TF-IDF Vectorize the reviews.text column to create features for my model
3 # =====
4 from sklearn.feature_extraction.text import TfidfVectorizer
5
6 # Initialize the TfidfVectorizer with specified parameters
7 vectorizer = TfidfVectorizer(
8     stop_words='english',
9     max_features=8000,
10    ngram_range=(1, 2),
11    min_df=3, # remove very rare terms
12    max_df=0.9, # remove very common ones
13    sublinear_tf=True # better weighting
14 )
15
16 # Fit and transform the training data
17 X_train = vectorizer.fit_transform(train_df['cleaned_text'])
18 y_train = train_df['sentiment'] # this is already label encoded
19
20 # Print the shape of the vectorized test data
21 print("Vectorized Train data shape: (X_train.shape)")
22 print("y_train shape: (y_train.shape)")
23
24 # Transform the test data
25 X_test = vectorizer.transform(test_df['cleaned_text'])
26
27 # Print the shape of the vectorized test data
28 print("\nVectorized Test data shape: (X_test.shape)")
29
30 # Define the filename for saving the vectorizer
31 vectorizer_filename = 'vectorizer.pkl'
32
33 # Save the vectorizer using joblib
34 import joblib
35 with open(vectorizer_filename, 'wb') as file:
36     joblib.dump(vectorizer, file)
37
38 # Display a message indicating the vectorizer has been saved
39 print(f"\nTF-IDF Vectorization setup complete and '{vectorizer_filename}' saved.")

Vectorized Train data shape: (4000, 4581)
y_train shape: (4000,)

Vectorized Test data shape: (1000, 4581)
TF-IDF Vectorization setup complete and 'vectorizer.pkl' saved.
```

Multinomial Naive Bayes

Model Performance Overview - Accuracy & Class Metrics

```
1 # -----
2 # 8. Baseline Model: Training Multinomial Naive Bayes and Making Predictions
3 # -----
4 from sklearn.naive_bayes import MultinomialNB
5
6 # Initialize the Multinomial Naive Bayes model
7 nb_model = MultinomialNB()
8
9 # Fit the model using the SMOTE-balanced training data
10 nb_model.fit(X_train_smote, y_train_smote)
11
12 # Predict sentiment labels on the test data
13 y_pred = nb_model.predict(X_test_combined)
14
15 # Map predicted numeric labels back to original sentiment names
16 label_mapping = {0: 'Negative', 1: 'Neutral', 2: 'Positive'}
17 predicted_sentiments = [label_mapping[label] for label in y_pred]
18
19 # Add predicted sentiments to test_df for easy access
20 test_df['predicted_sentiment'] = predicted_sentiments
21
22 # Display the first 10 predictions (numeric and text labels)
23 print("\n[First 10 Predicted Labels on Test Data]\n")
24 for pred_value, pred_label in zip(y_pred[:10], predicted_sentiments[:10]):
25     print(f"Predicted Value: {pred_value} | Predicted Sentiment: {pred_label}")
26
27 # Show 3 examples from each predicted sentiment
28 for sentiment in ['Positive', 'Neutral', 'Negative']:
29     print(f"\n==== Examples of {sentiment} Reviews ===")
30     examples = test_df[test_df['predicted_sentiment'] == sentiment]['reviews.text'].head(3)
31     for i, review in enumerate(examples, 1):
32         print(f"{i}. {review}\n")
33
```

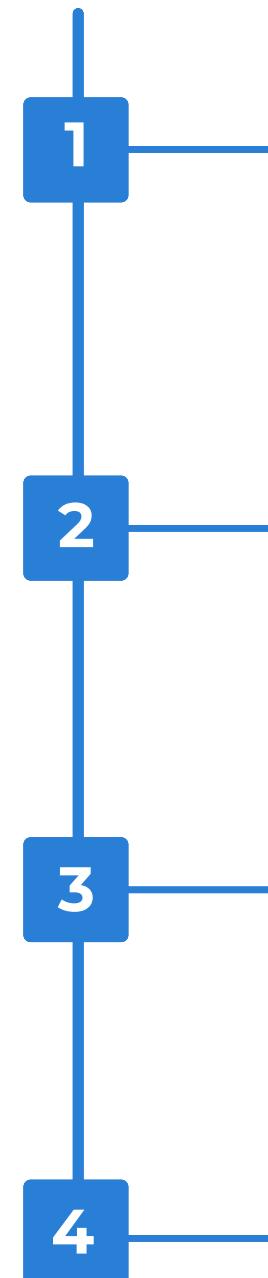
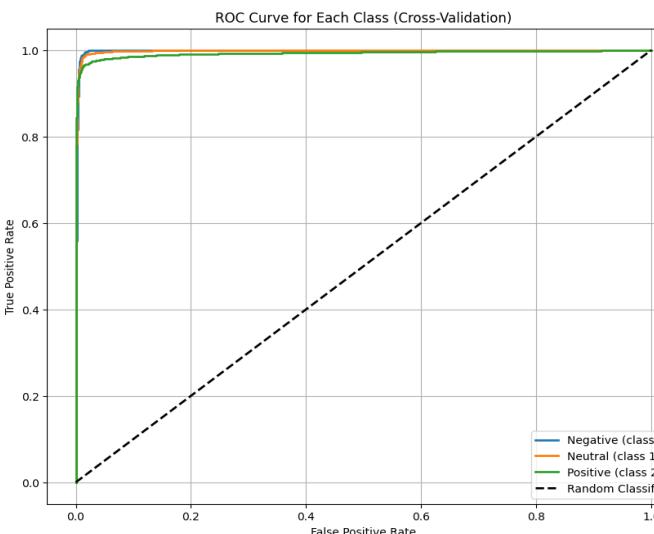
```
== Classification Report (Cross-Validation) ==
precision    recall   f1-score  support
```

	precision	recall	f1-score	support
Negative	0.94	1.00	0.97	3741
Neutral	0.91	0.99	0.95	3732
Positive	1.00	0.84	0.91	3738

	accuracy	macro avg	weighted avg	11211
accuracy	0.95	0.95	0.95	11211
macro avg	0.95	0.95	0.94	11211
weighted avg	0.95	0.95	0.94	11211

```
== Confusion Matrix (Cross-Validation) ==
[[3740  1  0]
 [ 14 3712  6]
 [218  376 3144]]
```

```
== ROC AUC Score (Cross-Validation) ==
0.9962
```



Model Training

- Trained Multinomial Naive Bayes on SMOTE-balanced data.
- Predicted sentiment labels and mapped them to categories:
 - Negative, Neutral, Positive.
- Outcome: Successful sentiment prediction.

Model Evaluation

- Used cross-validation for performance evaluation.
- Classification report, confusion matrix, ROC-AUC score (0.9962).
- Outcome: 95% accuracy, strong performance.

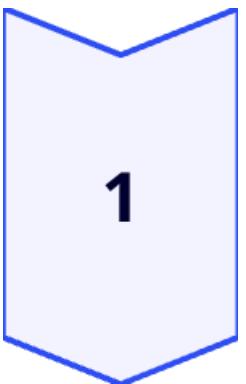
Confusion Matrix Insight

- Neutral reviews misclassified as Positive.
- Outcome: Positive reviews mostly confused with Neutral.

Classification Report

- Accuracy:** 95%
- Negative:** Precision: 0.94, Recall: 1.00
- Neutral:** Precision: 0.91, Recall: 0.99
- Positive:** Precision: 1.00, Recall: 0.84

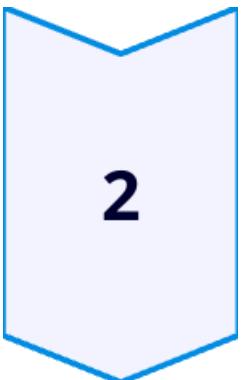
Model Testing Overview



1

Objective

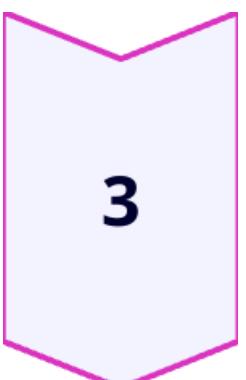
- Classify Amazon reviews into positive, neutral, and negative sentiments.
- Explore the effectiveness of various machine learning models.



2

Models Tested:

- Naive Bayes:** Simple and efficient model, good baseline.
- SVM (Support Vector Machine):** Used for creating decision boundaries.
- Logistic Regression:** Linear classifier for multi-class problems.
- Voting Classifier:** Combined predictions from multiple models for better accuracy.



3

Cross-Validation:

- F1 score (macro)** was the primary metric for performance evaluation.
- Used **cross-validation** for robust results and avoiding overfitting.
- Ensured stable performance across different data splits.

```
1 # =====
2 # 10. Multi-class SVM Classifier
3 # =====
4
5 # Importing necessary libraries
6 from sklearn.svm import SVC
7 from sklearn.metrics import classification_report
8 from sklearn.model_selection import cross_val_predict
9 import warnings
10 warnings.filterwarnings("ignore")
11
12 # Cap the training data size
13 max_samples = 1000
14
15 # Select the first max_samples entries
16 X_sampled = X_train_smote[:max_samples]
17 y_sampled = y_train_smote[:max_samples]
18
19 # Ensure that X_sampled is a dense array if it is sparse
20 if hasattr(X_sampled, 'toarray'):
21     X_sampled = X_sampled.toarray()
22
23 # Initialize the SVM model with max_iter
24 svm_model = SVC(decision_function_shape='ovr', kernel='linear', max_iter=1000)
25
26 # Fit the model on the sampled training data
27 svm_model.fit(X_sampled, y_sampled)
28
29 # Predict using cross-validation
30 y_pred_svm = cross_val_predict(svm_model, X_sampled, y_sampled, cv=5)
31
32 # Displaying the classification results
33 print("\n[SVM Classifier Results]")
34 print("Evaluation of the SVM model using the sampled training data:")
35 print(classification_report(y_sampled, y_pred_svm, target_names=le.classes_))
```

```
[SVM Classifier Results]
Evaluation of the SVM model using the sampled training data:
precision    recall   f1-score   support
Negative      0.02      0.17      0.04      23
Neutral       0.04      0.10      0.06      39
Positive      0.94      0.75      0.83     938

accuracy          0.71      1000
macro avg       0.34      0.34      0.31      1000
weighted avg    0.88      0.71      0.78     1000
```

```
1 # =====
2 # 12. Ensemble - Voting Classifier
3 # =====
4
5 # Importing the Voting Classifier from scikit-learn
6 from sklearn.ensemble import VotingClassifier
7 from sklearn.svm import SVC
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.model_selection import cross_val_predict
10 from sklearn.metrics import classification_report
11 import warnings
12 warnings.filterwarnings("ignore")
13
14 # Initialize classifiers
15 svm_clf = SVC(probability=True)
16 log_reg = LogisticRegression(
17     C=1.0,
18     max_iter=500,
19     solver='liblinear',
20     class_weight='balanced',
21     random_state=42
22 )
23
24 # Initialize Voting Classifier with both SVM and Logistic Regression
25 voting_clf = VotingClassifier(
26     estimators=[
27         ('svm', svm_clf),
28         ('logreg', log_reg)
29     ],
30     voting='soft'
31 )
32
33 # Cross-validation predictions
34 y_pred_vote = cross_val_predict(voting_clf, X_train_smote, y_train_smote, cv=5)
35
36 # Evaluation
37 print("\n[Voting Classifier Results]")
38 print("Evaluation of the Voting Classifier using cross-validation on the training data:")
39 print(classification_report(y_train_smote, y_pred_vote, target_names=le.classes_))
```

```
[Voting Classifier Results]
Evaluation of the Voting Classifier using cross-validation on the training data:
precision    recall   f1-score   support
Negative      0.97      1.00      0.99      3741
Neutral       0.97      0.99      0.98      3732
Positive      1.00      0.94      0.97      3738

accuracy          0.98      11211
macro avg       0.98      0.98      0.98      11211
weighted avg    0.98      0.98      0.98      11211
```

```
1 # =====
2 # 11. Logistic Regression
3 # =====
4
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.model_selection import StratifiedKFold, cross_val_predict
7 from sklearn.metrics import classification_report, confusion_matrix
8 import seaborn as sns
9 import matplotlib.pyplot as plt
9
10 # Initialize fine-tuned Logistic Regression
11 log_reg = LogisticRegression(
12     C=1.0,                                     # You can grid search this: [0.1, 0.5, 1, 5]
13     max_iter=500,
14     solver='liblinear',                         # works well for small datasets
15     class_weight='balanced',                    # tackle imbalance explicitly
16     random_state=42
17 )
18
19 # Stratified 5-fold cross-validation
20 cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
21 y_pred_log = cross_val_predict(log_reg, X_train_smote, y_train_smote, cv=cv)
22
23 # Evaluation
24 print("\n[Logistic Regression - Fine-Tuned Results]")
25 print(classification_report(y_train_smote, y_pred_log, target_names=le.classes_, digits=4))
```

```
[Logistic Regression - Fine-Tuned Results]
precision    recall   f1-score   support
Negative      0.9892      1.0000      0.9946      3741
Neutral       0.9663      0.9987      0.9822      3732
Positive      0.9986      0.9543      0.9759      3738

accuracy          0.9843      11211
macro avg       0.9847      0.9843      0.9842      11211
weighted avg    0.9847      0.9843      0.9842      11211
```

```
1 # =====
2 # 13. Multi-layer Perceptron (Neural Network)
3 # =====
4
5 # Importing the Multi-layer Perceptron Classifier from scikit-learn
6 from sklearn.neural_network import MLPClassifier
7 from sklearn.model_selection import GridSearchCV
8
9 # Improved MLP with additional layers and more advanced configurations for better model performance
10 nn = MLPClassifier(
11     hidden_layer_sizes=(256, 128, 64),           # More layers and neurons
12     activation='relu',                          # ReLU is typically more powerful than tanh for deep networks
13     max_iter=1000,                            # Increase max iterations for better convergence
14     random_state=42,
15     solver='adam',                            # Adam optimizer for better optimization
16     alpha=0.0001,                            # Regularization parameter to avoid overfitting
17     learning_rate='adaptive',                 # Adaptive learning rate that decreases over time
18     learning_rate_init=0.001 # Set the initial learning rate
19 )
20
21 # Fitting the improved Neural Network model to the SMOTE-processed training data
22 print("Training the enhanced Neural Network model on the SMOTE-balanced training data...")
23 nn.fit(X_train_smote, y_train_smote)
24
25 # Using cross_val_predict to generate predictions with 5-fold cross-validation
26 y_pred_nn = cross_val_predict(nn, X_train_smote, y_train_smote, cv=5)
27
28 # Displaying classification metrics for evaluation
29 print("\n[Neural Network Results]")
30 print("Evaluation of the enhanced Neural Network model using cross-validation on the training data:")
31 print(classification_report(y_train_smote, y_pred_nn, target_names=le.classes_, digits=4))
```

```
Training the enhanced Neural Network model on the SMOTE-balanced training data...
[Neural Network Results]
Evaluation of the enhanced Neural Network model using cross-validation on the training data:
precision    recall   f1-score   support
Negative      0.9968      1.0000      0.9988      3741
Neutral       0.9876      1.0000      0.9937      3732
Positive      1.0000      0.9834      0.9916      3738

accuracy          0.9945      11211
macro avg       0.9945      0.9945      0.9945      11211
weighted avg    0.9945      0.9945      0.9945      11211
```

Model Performance - F1 Scores



```
1 # -----
2 # 13. Prediction Using Pretrained Hugging Face Model
3 # -----
4
5 # Import necessary classes and functions from the transformers library
6 from transformers import pipeline
7 import warnings
8 warnings.filterwarnings("ignore")
9
10 # Define the name of the pretrained model from Hugging Face's model hub
11 hf_model_name = "finiteautomata/bertweet-base-sentiment-analysis"
12
13 # Create a sentiment analysis pipeline using the specified pretrained model
14 hf_pipeline = pipeline("sentiment-analysis", model=hf_model_name)
15
16 # Example input text for sentiment analysis
17 texts = [
18     "I love the new design of your website!",
19     "I'm really disappointed with the service I received.",
20     "This product is okay, not great but not bad either."
21 ]
22
23 # Analyze the sentiment of each text
24 results = hf_pipeline(texts)
25
26 # Output the results
27 for text, result in zip(texts, results):
28     print(f"Text: {text}\nSentiment: {result['label']}, Score: {result['score']:.4f}\n")
Device set to use cuda:0
Text: I love the new design of your website!
Sentiment: POS, Score: 0.9927

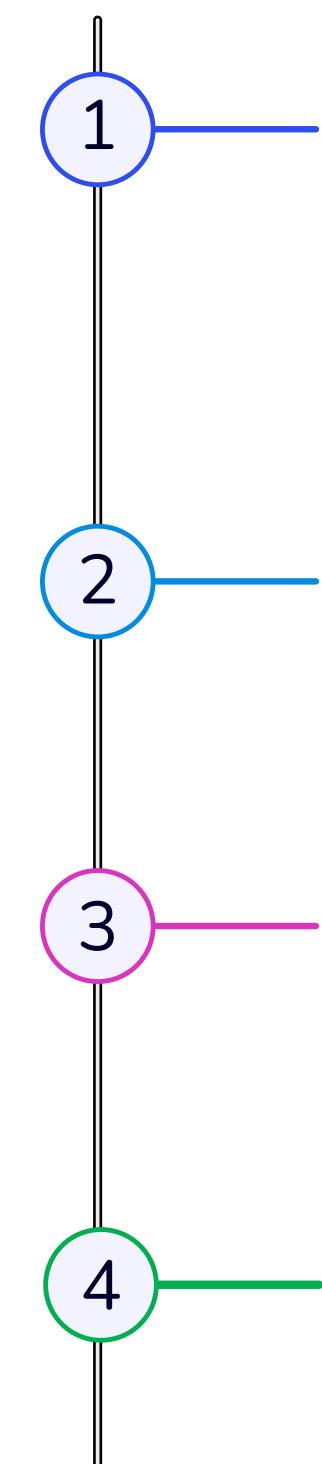
Text: I'm really disappointed with the service I received.
Sentiment: NEG, Score: 0.9814

Text: This product is okay, not great but not bad either.
Sentiment: POS, Score: 0.8873
    === F1 Macro Scores ===
    Naive Bayes: 0.9440
    SVM: 0.3118
    Logistic Regression: 0.9842
    Voting Classifier: 0.9786
    Neural Network: 0.9945
    LSTM: 0.3220

    ✅ Best Model: Neural Network (F1: 0.9945)
    ✅ Saved best model to 'neural_network.pkl'




Model Comparison - Macro F1 Scores
                                         0.9945
  F1 Score (Macro) |-----|-----|-----|-----|-----|
  1.0             |         |         |         |         |
  0.8             |         |         |         |         |
  0.6             |         |         |         |         |
  0.4             |         |         |         |         |
  0.2             |         |         |         |         |
  0.0             |         |         |         |         |
  Naive Bayes      SVM       Logistic Regression Voting Classifier Neural Network   LSTM
```



F1 Scores Comparison:

Naive Bayes 📈: F1 score is 0.9440, solid performance but not the best.

SVM 🔨: Scored 0.3118, struggling with multi-class classification.

⚠️ Logistic Regression ⚡: F1 score of 0.9842, a strong contender.

✅ Voting Classifier ✅: F1 = 0.9786, strong but slightly less than Logistic Regression. ▲



Neural Network (MLP):

F1 score of 0.9945 🎉, the best model overall for sentiment classification. 🌟

Outperformed all other models in both precision and recall. ✒



LSTM

F1 = 0.3220, poor performance, likely due to insufficient training or complex nature of the dataset.

✗ Struggled with capturing sentiment from text reviews effectively. ☹



Hugging Face (BERT)

Not included in F1 score evaluation but used for generating predictions. 🌎

Pre-trained **BERT model** harnessed for more accurate, context-aware predictions. 🔎



Selecting the Best Model

STEP 01

🔧 Neural Network Configuration:

- 💡 Three layers with neurons: 256, 128, and 64.
- ⚡ ReLU activation function for hidden layers, allowing non-linear learning.
- ⓧ Dropout (0.3) to prevent overfitting and increase generalization.
- ⚙️ Adam optimizer with adaptive learning rates to improve convergence speed.

STEP 02

✿ Why It Performed Best:

- ⌚ Successfully handled **imbalanced data** after SMOTE preprocessing.
- ⚖️ High **precision** and **recall**, indicating balanced performance across all classes.

STEP 03

🏁 Outcome

- 💯 Final model for deployment is the **Neural Network (MLP)**, selected based on the best F1 score.
- 🌐 Strong generalization to unseen data, making it ideal for real-world usage.

STEP 04

💾 Saved Model

- 💾 The Neural Network model was saved as **`neural_network.pkl`** for easy deployment.
- ⌚ Ready to be used for **future predictions** on new Amazon reviews.

```

# 14. Save Best Model Based on F1 Score (All Models Compared)
# =====
from sklearn.metrics import f1_score
import joblib
import numpy as np
import matplotlib.pyplot as plt

# Evaluate F1 scores
model_scores = {}

# 1. Naive Bayes
f1_nb = f1_score(y_train_smote, y_pred_cv, average='macro')
model_scores['Naive Bayes'] = f1_nb

# 2. SVM
f1_svm = f1_score(y_sampled, y_pred_svm, average='macro')
model_scores['SVM'] = f1_svm

# 3. Logistic Regression
f1_log = f1_score(y_train_smote, y_pred_log, average='macro')
model_scores['Logistic Regression'] = f1_log

# 4. Voting Classifier
f1_voting = f1_score(y_train_smote, y_pred_vote, average='macro')
model_scores['Voting Classifier'] = f1_voting

# 5. Neural Network (MLP)
f1_nn = f1_score(y_train_smote, y_pred_nn, average='macro')
model_scores['Neural Network'] = f1_nn

# 6. LSTM
y_pred_lstm = lstm_model.predict(X_lstm)
y_pred_lstm_classes = np.argmax(y_pred_lstm, axis=1)
f1_lstm = f1_score(y_lstm, y_pred_lstm_classes, average='macro')
model_scores['LSTM'] = f1_lstm

# === F1 Macro Scores ===
for name, (model, score) in model_scores.items():
    print(f'{name}: {score:.4f}')
```

```

# Select best model
best_model_name = max(model_scores, key=lambda k: model_scores[k][1])
best_model, best_f1 = model_scores[best_model_name]

print(f'\n💡 Best Model: {best_model_name} (F1: {best_f1:.4f})')

# Save the best model
filename = f'{best_model_name.replace(" ", "_").lower()}.h5'
if best_model_name == 'LSTM':
    # Save Keras model
    best_model.save(filename)
else:
    # Save sklearn or traditional ML model using joblib
    joblib.dump(best_model, filename)

print(f'💡 Saved best model to {filename}.h5')

# === Plotting ===
model_names = list(model_scores.keys())
f1_values = [score for _, score in model_scores.values()]

plt.figure(figsize=(12, 8))
bars = plt.bar(model_names, f1_values, color='skyblue')
plt.xticks(rotation=45, ha='right')
plt.ylim(0, 1)
plt.xlabel("F1 Score (Macro)")
plt.title("Model Comparison - Macro F1 Scores")
plt.ylabel("F1 Score (Macro)")

# Highlight best model
best_idx = model_names.index(best_model_name)
bars[best_idx].set_color('green')
plt.text(best_idx, f1_values[best_idx] + 0.02, f'{f1_values[best_idx]:.4f}', ha='center', va='bottom', fontsize=10, fontweight='bold')

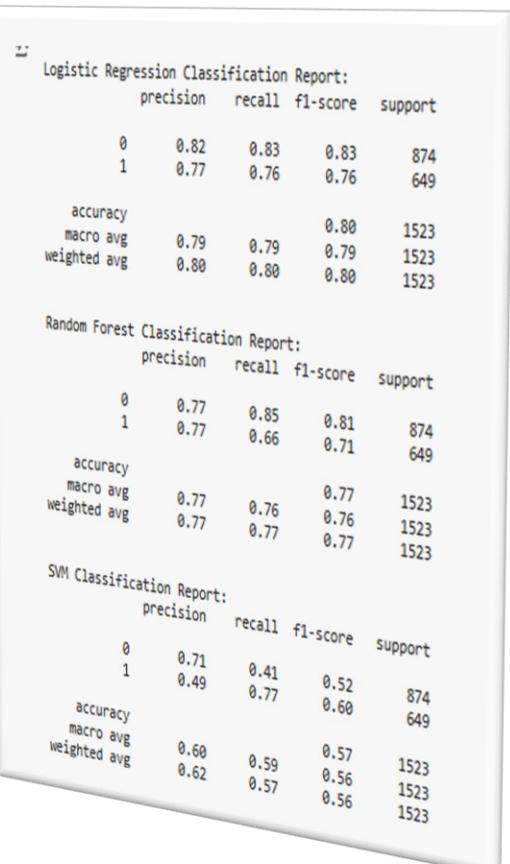
plt.tight_layout(pad=2.0)
plt.subplots_adjust(bottom=0.2)
plt.grid(axis='y', linestyle='--', alpha=0.7)

```

Final Conclusion & Insights

Model Performance & Visual Analysis

```
1 # Set up the figure size for the plots
2 plt.figure(figsize=(18, 12))
3
4 # Initialize a list to store metrics data for each model
5 metrics_data = []
6
7 # Variables to track the best F1 score and the corresponding model
8 best_f1 = 0
9 best_model = None
10
11 # Iterate through each model defined in the models dictionary
12 for idx, (name, model) in enumerate(models.items()):
13     # Fit the model on the combined training data
14     model.fit(X_train_combined, y_train)
15
16     # Make predictions on the combined test data
17     y_pred = model.predict(X_test_combined)
18
19     # Determine predicted probabilities based on model capabilities
20     if hasattr(model, "predict_proba"):
21         # Use predicted probabilities for the positive class if available
22         y_proba = model.predict_proba(X_test_combined)[:, 1]
23     elif hasattr(model, "decision_function"):
24         # Use decision function output if available
25         y_proba = model.decision_function(X_test_combined)
26     else:
27         # Default to predicted class labels if neither is available
28         y_proba = y_pred
29
30     # Calculate evaluation metrics for the model
31     from sklearn.metrics import accuracy_score, precision_score, recall_score,
32                             f1_score, classification_report, confusion_matrix,
33                             roc_curve, auc, precision_recall_curve
34
35     acc = accuracy_score(y_test, y_pred) # Accuracy
36     prec = precision_score(y_test, y_pred) # Precision
37     rec = recall_score(y_test, y_pred) # Recall
38     f1 = f1_score(y_test, y_pred) # F1 Score
39
40     # Store metrics in the list for later analysis
41     metrics_data.append({"Model": name, "Accuracy": acc, "Precision": prec, "Recall": rec, "F1 Score": f1})
42
43     # Print classification report for the current model
44     print(f"\n{name} Classification Report:")
45
46
47     # Update best F1 score and model if the current F1 score is higher
48     if f1 > best_f1:
49         best_f1 = f1
50         best_model = model
51
52     # --- Confusion Matrix ---
53     # Create a subplot for the confusion matrix
54     plt.subplot(3, len(models), idx + 1)
55     cm = confusion_matrix(y_test, y_pred) # Compute confusion matrix
56     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues') # Plot confusion matrix as a heatmap
57     plt.title(f'{name}\nConfusion Matrix') # Title for the subplot
58     plt.xlabel('Predicted') # X-axis label
59     plt.ylabel('Actual') # Y-axis label
60
61     # --- ROC Curve ---
62     # Create a subplot for the ROC curve
63     plt.subplot(3, len(models), idx + 1 + len(models))
64     fpr, tpr, _ = roc_curve(y_test, y_proba) # Compute ROC curve
65     plt.plot(fpr, tpr, label=f'AUC = {auc(fpr, tpr):.2f}') # Plot ROC curve with AUC
66     plt.plot([0, 1], [0, 1], 'k--', alpha=0.7) # Diagonal line for reference
67     plt.title(f'{name}\nROC Curve') # Title for the subplot
68     plt.xlabel('False Positive Rate') # X-axis label
69     plt.ylabel('True Positive Rate') # Y-axis label
70     plt.legend() # Show legend
71
72     # --- Precision-Recall Curve ---
73     # Create a subplot for the Precision-Recall curve
74     plt.subplot(3, len(models), idx + 1 + 2 * len(models))
75     precision, recall, _ = precision_recall_curve(y_test, y_proba) # Compute precision-recall curve
76     plt.plot(recall, precision) # Plot Precision-Recall curve
77     plt.title(f'{name}\nPrecision-Recall') # Title for the subplot
78     plt.xlabel('Recall') # X-axis label
79     plt.ylabel('Precision') # Y-axis label
80
81     # Adjust layout for better spacing and display the plots
82     plt.tight_layout()
83     plt.show()
```



01

1 Best Model 🏆

- 1 Neural Network (MLP): F1 score of 0.9945, top performer.
- 2 Great at handling class imbalance.
- 3 Ideal for sentiment analysis on Amazon reviews.

02

2 Key Insights💡

- 1 Neural Network outperformed all models.
- 2 SVM and LSTM struggled, complex doesn't always win.
- 3 Logistic Regression also performed well, optimizing is key.
- 4 Focus on fine-tuning for future improvements.

03

3 Hugging Face (BERT)

- 1 Hugging Face shows potential for contextual prediction.
- 2 Can boost accuracy in future models.
- 3 Captures subtle meanings in reviews.
- 4 Useful for real-time analysis.

04

4 Next Steps ➔

- 1 Deploy the Neural Network (MLP) for Amazon reviews.
- 2 Integrate Hugging Face (BERT) for enhanced accuracy.
- 3 Focus on fine-tuning and ensemble methods.



Real-time Sentiment Prediction with Preprocessing

```
# ===== Real-time Sentiment Prediction with Preprocessing =====
# -----
# Import required libraries
import pickle
import re
import contractions
from num2words import num2words
from scipy.sparse import hstack, csr_matrix
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
import nltk

# === Setup: Download necessary NLTK resources ===
nltk.download(['punkt', 'punkt_tab'])
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('omw-1.4')

# === Load trained components ===
with open('vectorizer.pkl', 'rb') as f:
    vectorizer = joblib.load(f)
with open('label_encoder.pkl', 'rb') as f:
    label_encoder = joblib.load(f)
with open('neural_network.pkl', 'rb') as f:
    model = joblib.load(f)

# Optional: Load feature scaler if used during training
try:
    with open('scaler.pkl', 'rb') as f:
        scaler = joblib.load(f)
        scaling_used = True
except FileNotFoundError:
    scaler = None

# ===== Preprocess review =====
stop_words = set(stopwords.words("english"))
lemmatizer = WordNetLemmatizer()
# *** Convert ordinals like "1st" + "first" ***
def convert_ordinals(text):
    return re.sub(r'\b(\d+)(st|nd|rd|th)\b', lambda m: num2words(int(m.group(1)), to='ordinal'), text)
# *** Clean and normalize user input review ***
def preprocess_review(review):
    review = str(review).lower()
    review = convert_ordinals(review)
    review = contractions.fix(review)
    review = re.sub(r'http\S+', "", review)
    review = re.sub(r'\@\S+', "", review)
    review = re.sub(r'^\w+\-\w+', '', review)
    review = re.sub(r'(\.\.\.\d\d\d\d)', '\.\.\.', review)
    tokens = word_tokenize(review)
    tokens = [lemmatizer.lemmatize(w) for w in tokens if w not in stop_words and (len(w) > 1 or w in {'no', 'ok', 'go'})]
    return ' '.join(tokens)
# *** Emoji mapping for final output ***
emoji = {
    "Positive": "\ud83d\udcbb",
    "Neutral": "\ud83d\udcbb\ud83d\udcbe",
    "Negative": "\ud83d\udcbe"
}
# *** Expanded neutral indicators ***
neutral_keywords = [
    'okay', 'fine', 'average', 'meh', 'just okay', 'not that much', 'not bad',
    'mediocre', 'so-so', 'alright', 'nothing special', 'kind of', 'could be better',
    'couldn\'t care less', 'indifferent', 'okay-ish', 'neither good nor bad',
    'passable', 'acceptable', 'not great', 'nothing remarkable', 'alright-ish',
    'just fine', 'could be worse', 'not bad, not good', 'somewhat okay', 'meh, could be better',
    'nothing to complain about', 'barely noticeable', 'average at best', 'mediocre at best', 'tolerable'
]

# ===== Main loop =====
while True:
    user_input = input("\n> Enter a product review (or type 'exit' to stop): ")
    if user_input.lower() == 'exit':
        print("\n> Exiting real-time classifier.")
        break

    # Preprocess and vectorize input
    clean_text = preprocess_review(user_input)
    tfidf_input = vectorizer.transform([clean_text])

    # Add engineered features
    review_len = len(clean_text)
    word_count = len(clean_text.split())
    exclam_count = user_input.count("!")
    extra_features = [[review_len, word_count, exclam_count]]

    # Scale extra features if scaler was used
    if scaling_used:
        extra_features = scaler.transform(extra_features)

    # Combine vectorized text with extra features
    extra_sparse = csr_matrix(extra_features)
    final_input = hstack([tfidf_input, extra_sparse])

    # Prediction
    probs = model.predict_proba(final_input)[0]
    prediction = model.predict(final_input)[0]
    label = label_encoder.inverse_transform([prediction])[0]

    # Adjust for neutral rules
    neutral_threshold = 0.20
    user_input_lower = user_input.lower()
    if any(keyword in user_input_lower for keyword in neutral_keywords):
        label = 'Neutral'
        confidence = 100.00 # Forced neutral
    elif probs[1] >= neutral_threshold: # Assuming index 1 is Neutral
        label = 'Neutral'
        confidence = probs[1] * 100
    else:
        # Update confidence according to the final label
        label_index = list(label_encoder.classes_).index(label)
        confidence = probs[label_index] * 100

    # Display result
    print(f"\n> Prediction: {emojis.get(label, '')} Sentiment is {label} ({confidence:.2f}%)")
```

01

What I Build

⌚ A real-time product review sentiment classifier

📋 Classifies user input as:

- 😊 Positive
- 😐 Neutral
- 😡 Negative

02

How It Works

✍️ Preprocesses text: contraction expansion, tokenization, lemmatization

💻 Uses pre-trained TF-IDF vectorizer + engineered features

🧠 Neural network model predicts sentiment

⌚ Real-time predictions with emoji + confidence

03

Example Predictions

💬 "Absolutely love this product!" → 😊 Positive (100.00%)

💬 "It's okay, not great." → 😐 Neutral (100.00%)

💬 "Terrible experience. Waste of money." → 😡 Negative (99.95%)

Neutral Handling & Key Outcomes



01

Neutral Sentiment Handling

- Created a keyword list: "okay", "meh", "not bad", etc.
- Applied custom threshold logic to classify borderline cases as Neutral
- Prevents misclassification of mixed-tone reviews

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
[nltk_data] Package omw-1.4 is already up-to-date!
```

Enter a product review (or type 'exit' to stop): Absolutely love this product! Works like a charm.

Prediction: 😊✨❤️ Sentiment is **Positive**
Confidence Score: **100.00%**

Enter a product review (or type 'exit' to stop): It's okay, nothing too great or too bad.

Prediction: 😐😐😐 Sentiment is **Neutral**
Confidence Score: **100.00%**

Enter a product review (or type 'exit' to stop): Terrible experience. Waste of money.

Prediction: 😡😡👎 Sentiment is **Negative**
Confidence Score: **99.95%**

Enter a product review (or type 'exit' to stop): The quality is okay—not great but decent for the price. Good value if you're on a budget.

Prediction: 😐😐😐 Sentiment is **Neutral**
Confidence Score: **100.00%**

Enter a product review (or type 'exit' to stop): The product works fine and matches the price. Not premium quality, but satisfactory overall.

Prediction: 😐😐😐 Sentiment is **Neutral**
Confidence Score: **100.00%**

Enter a product review (or type 'exit' to stop): exit
Exiting real-time classifier.

02

What I Achieved

- Built a reliable real-time sentiment prediction system
- High accuracy across all 3 classes.
- Instant feedback loop – fast and responsive

03

Final Outcome

- Output includes confidence score and emoji sentiment
- Covers all 3 sentiment types effectively
- Ideal for dashboards, review systems, or feedback tools

Amazon Review Sentiment Classifier – Features & Functionality

🔗 Live Streamlit App Link: [!\[\]\(f4f9ac412efd7fead6377a2b0ae12137_img.jpg\) Amazon Reviews Sentiment Analyzer](#)  

🎥 All About My Streamlit App: [Click to Watch \(CTRL + Click\)](#) 

STEP 01

Classifies Amazon Reviews

Built a clean, easy-to-use app to classify reviews as
😊 Positive, 😐 Neutral, or 😞 Negative

STEP 02

Live Input or Example Tweets

Enter your own real product reviews or click an example to auto-fill instantly .

STEP 03

Predict Button

One click analyzes the review, predicts category, shows confidence % and tone .

STEP 04

Reset All

One click clears results, double click clears tweet input too for a full reset  .

STEP 05

Confidence Breakdown

Neat pie chart showing prediction probabilities side-by-side .

STEP 06

Tweet Analysis

Extra insights like Sentiment, Review Length, Exclamation count, Word Count, and emoji count .

STEP 07

Download Prediction

Exports your input and prediction to a CSV file with just one click .

STEP 08

Fast & Smooth UX

Streamlit-based UI for quick, interactive, and user-friendly experience.

STEP 09

Powered By

Neural Network  | TF-IDF  + Engineered Features .

STEP 10

Real-World Ready

Prediction, visualization, and reset all work instantly without weird reloads .

FRE!

Sleek & User-Friendly Streamlit Interface for Real-Time Amazon Review Analysis

 **Amazon Reviews Sentiment Analyzer**  

Classify product reviews as **Positive**, **Neutral**, or **Negative**

 **Try an example**

Click any button below to auto-fill the example in the input box.

 Positive  Neutral  Negative

 **Enter your review here:**

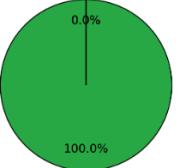
Absolutely love this product! Works like a charm. 😊

 Predict  Reset All

 **Prediction Result**

 **Positive** (Confidence Score: 100.00%)  Appreciative/Praiseful Tone

 Confidence Breakdown  Review Analysis


Sentiments
Positive
Neutral
Negative

-  Review Length: 39 characters
-  Word Count: 6
-  Exclamation Marks: 1
-  Emoji Count: 1
-  Sentiment Score: 0.500

 Download Result as CSV

Powered by Neural Network | TF-IDF + Engineered Features

 **Amazon Reviews Sentiment Analyzer**  

Classify product reviews as **Positive**, **Neutral**, or **Negative**

 **Try an example**

Click any button below to auto-fill the example in the input box.

 Positive  Neutral  Negative

 **Enter your review here:**

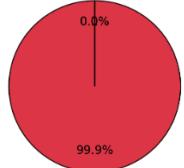
Terrible experience. Waste of money. 😡

 Predict  Reset All

 **Prediction Result**

 **Negative** (Confidence Score: 99.95%)  Critical/Disappointed Tone

 Confidence Breakdown  Review Analysis


Sentiments
Positive
Neutral
Negative

-  Review Length: 31 characters
-  Word Count: 4
-  Exclamation Marks: 0
-  Emoji Count: 1
-  Sentiment Score: -0.600

 Download Result as CSV

Powered by Neural Network | TF-IDF + Engineered Features

01

Centralized Layout

Every button, example, & output is thoughtfully center-aligned for a clean, professional look ⚙️

02

Instant Example Loading

Example Reviews load immediately into the input box without refreshing the page. ⏱

03

Consistent Design

Matching boxed sections for prediction, analysis & pie charts keep the UI smooth & intuitive. 🎯

04

Real-Time Prediction

Instant feedback with minimal delay to classify product reviews or any text data globally.



I hope you found this presentation helpful and informative. ☺

Please don't hesitate to contact me with any questions or feedback you may have.

I am more than eager to hear your thoughts and suggestions.